

Data Engineering Bootcamp Report

This document has the answers for questions 1 through 6 of the data bootcamp. For the script I decided to use OOP since I find it quite comfortable for these types of analysis. I created a method to read in the file every time a method is run. This is because when working with big files in my current job I have to be really efficient with memory use. I do not have the luxury of saving the file in more dataframes.

Question 1: How many commercial chains are monitored, and therefore, included in this database?

```
# question 1 answer
def chains_cnt(self):
    """
    Prints the total number of unique chains.

    Reads the 'cadenaComercial' column to obtain all the existing
    commercial chains without duplicates.

    Returns:
        None

    Typical usage example:
    analysis = ExploratoryAnalysis(file_name)
    analysis.chains_cnt()
    """
    self._read_columns('cadenaComercial')
    print(self.df.nunique())
    del self.df
    gc.collect()
```

For this question I used pandas' "nunique" method which basically returns the total sum of unique values of "cadenaComercial". This comes after reading in the chain column which is cleaned by removing any null values. I encountered a few rows that had the header instead of any useful info so I removed those as well.

Answer:

704

Question 2: What are the top 10 monitored products by State?

```
# question 2 answer
def prod_by_state(self, top_products):
    """
    Prints the top n (by amount) products by state.

    Uses the state and product columns to group them together. Results
    in a dataframe with three columns: state, product and frequency.

    Args:
        top_products: int.
            Number of products to be printed per state.

    Returns:
        None

    Typical usage example:
    analysis = ExploratoryAnalysis(file_name)
    analysis.prod_by_state(10)
    """
    self._read_columns(['producto', 'estado'])

    # reset_index is used to convert a Series object (returned by .size())
    # to a DataFrame.
    # At the same time it keeps the state column as part of the dataframe
    # (which is the index in the Series)
    self.df = self.df.groupby(by=['estado', 'producto'], dropna=True).size()
    self.df = self.df.reset_index()

    # sort values in a descending fashion to get the top n products by state
    self.df.sort_values(by=['estado', 0],
                        ascending=[True, False],
                        inplace=True)
    self.df = self.df.groupby('estado').head(top_products)
    self.df.rename(columns={0: 'cantidad'}, inplace=True)
    self.df.reset_index(drop=True, inplace=True)

    print(self.df)
    del self.df
    gc.collect()
```

I decided to use “groupby” twice since I had to sort the values after the first use of “groupby” to get the top 10 products.

Partial answer to question 2:

	estado	producto	cantidad
0	AGUASCALIENTES	FUD	12005
1	AGUASCALIENTES	DETERGENTE P/ROPA	10188
2	AGUASCALIENTES	LECHE ULTRAPASTEURIZADA	9824
3	AGUASCALIENTES	SHAMPOO	9654
4	AGUASCALIENTES	REFRESCO	9481
5	AGUASCALIENTES	DESODORANTE	8859
6	AGUASCALIENTES	JABON DE TOCADOR	8517
7	AGUASCALIENTES	CHILES EN LATA	7946
8	AGUASCALIENTES	YOGHURT	7401
9	AGUASCALIENTES	MAYONESA	7173
10	BAJA CALIFORNIA	REFRESCO	37243
11	BAJA CALIFORNIA	DETERGENTE P/ROPA	23395
12	BAJA CALIFORNIA	FUD	19967
13	BAJA CALIFORNIA	SHAMPOO	19123
14	BAJA CALIFORNIA	JABON DE TOCADOR	18348
15	BAJA CALIFORNIA	CHILES EN LATA	16676
16	BAJA CALIFORNIA	GALLETAS	15873
17	BAJA CALIFORNIA	PANTALLAS	15703
18	BAJA CALIFORNIA	CEREALES	15398
19	BAJA CALIFORNIA	DESODORANTE	14748
20	BAJA CALIFORNIA SUR	REFRESCO	27770
21	BAJA CALIFORNIA SUR	FUD	17776

Question 3: Which is the commercial chain with the highest number of monitored products?

```
# question 3 answer
def top_chain(self):
    """
    Prints the chain with the most products.

    Reads the 'cadenaComercial' column to obtain the chain with the most
    products including the number of products for the chain.

    Returns:
        None

    Typical usage example:
    analysis = ExploratoryAnalysis(file_name)
    analysis.top_chain()
    """
    self._read_columns('cadenaComercial')
    self.df = self.df.value_counts().reset_index().iloc[0, :]
    self.df.rename({0: 'Numero de productos'}, inplace=True)
    print(self.df)
    del self.df
    gc.collect()
```

For this answer I used a value_counts which returns all the occurrences per each unique row of the given column. I used iloc just to get the first element since the question is only asking for the chain with the highest number of products. Fortunately for us pandas' value_counts method already returns each chain with their occurrences sorted.

Answer 3:

index	WAL-MART
Numero de productos	8643133

Question 4: Use the data to find an interesting fact.

```
Notes:
    We subtract 1e-6 to min_price only for the first price range since
    we want the inequality (self.df.values > min_price) to be inclusive
    only in the first step. Which is why we sum 1e-6 to max_price to
    keep it inclusive.
"""

self._read_float_columns(['fechaRegistro', 'precio'])

days_of_week = {0: 'Monday', 1: 'Tuesday', 2: 'Wednesday',
                 3: 'Thursday', 4: 'Friday', 5: 'Saturday',
                 6: 'Sunday'}

for day_idx in range(7):
    day = self.df['fechaRegistro'].dt.dayofweek.isin([day_idx])
    price = self.df[day]['precio']
    range_length = 10
    print(f'\nPrice Ranges for {days_of_week[day_idx]}\n')
    range_size = (price.max() - price.min()) / range_length
    for cnt, i in enumerate(range(range_length)):
        min_price = (price.min() + cnt * range_size)
        max_price = (min_price + 1e-6 + range_size)
        min_price = min_price if cnt > 0 else (min_price - 1e-6)
        print("Min: {:.2f}".format(min_price),
              'Max: {:.2f}'.format(max_price))

        cond = (price.values > min_price) & (price.values <= max_price)
        print(len(self.df[day][cond]))

del self.df
gc.collect()
```

I decided to check the price ranges per day of the week in the data. For each day there will be 10 price ranges printed out along with the number of prices per price range. Weekends are included as well.

Partial answer for question 4:

Price Ranges for Wednesday	Price Ranges for Tuesday
Min: 0.45 Max: 25500.32 13207672	Min: 0.60 Max: 30000.44 13151331
Min: 25500.32 Max: 51000.19 6467	Min: 30000.44 Max: 60000.28 2776
Min: 51000.19 Max: 76500.06 992	Min: 60000.28 Max: 90000.12 354
Min: 76500.06 Max: 101999.93 281	Min: 90000.12 Max: 119999.96 62
Min: 101999.93 Max: 127499.80 63	Min: 119999.96 Max: 149999.80 38
Min: 127499.80 Max: 152999.67 54	Min: 149999.80 Max: 179999.64 11
Min: 152999.67 Max: 178499.54 7	Min: 179999.64 Max: 209999.48 1
Min: 178499.54 Max: 203999.41 0	Min: 209999.48 Max: 239999.32 6
Min: 203999.41 Max: 229499.28 3	Min: 239999.32 Max: 269999.16 7
Min: 229499.28 Max: 254999.15 2	Min: 269999.16 Max: 299999.00 3

The days with the most products reported were Wednesday and Tuesday. What I found interesting is that each day's price ranges follow a Pareto distribution. The quantity of products changes but the shape of the distribution does not.

Question 5: What are the lessons learned from this exercise?

I learned that what I thought of Profeco's data was wrong. Since the number of products per price has something resembling a Pareto distribution instead of a Gaussian which was what I expected at first. Other than that, I learned that many simple products such as school notebooks appear in this database.

Question 6: Can you identify other ways to approach this problem? Explain.

To benefit Profeco one option I can think of is to try to predict the number of reports per day. This due to the fact that in the answer for question 4 I found that more reports were received on Wednesday and Tuesday. By predicting the number of products per day we can have three options to benefit Profeco:

1. Lower employees needed on days when the number of reports is predicted to be low. This way Profeco can offer a more flexible schedule for its employees.

Pros:

- Possible micro savings by not having as many employees available when not needed (assuming they get paid by the hour).

Cons:

- Prediction errors may cause a shortage of employees when more are needed.

2. Use automated recordings or email answers for reports which are common such as sodas.

Pros:

- Possible savings by automating part of the process. This could help ameliorate employee stress by having some people's reports solved by automated systems.
- People's satisfaction by possibly simplifying the process and potentially reducing wait times when demand is high for their services.

Cons:

- A system that works inefficiently would make the problem worse. People trying to make a report would probably just try to make the report in person.
3. Try to fix the root of the problem researching the causes for the reports on the days with the most appearances.

Pros:

- Possibly inexpensive alternative that could reduce the number of reports drastically.

Cons:

- Unknown variables which affect the number of occurrences.
- Root of the problem could be something Profeco is unable to change.